

From Prediction to Production

What Machine Learning Can and Cannot Do for Materials Discovery

Michigan Journal of Economics | Future Tech Platforms

The standard story about artificial intelligence and materials science runs as follows: machine learning models can now screen millions of candidate materials in the time it would take a human researcher to characterize a handful, collapsing a discovery process that once required decades of empirical trial into a computational exercise that runs overnight. The story is accurate but incomplete. Computational prediction and experimental production are separated by a gap that no algorithm currently bridges, and that gap has economic consequences for how quickly AI-assisted materials discovery can actually influence industrial processes, including the semiconductor manufacturing processes that underpin the hardware competition examined elsewhere in this issue.

This paper traces the materials discovery pipeline, identifies where machine learning generates genuine economic value, and explains where the bottlenecks that constrain the pace of realized discovery lie.

The Traditional Pipeline and Its Costs

Materials discovery prior to the machine learning era followed a largely empirical process. A chemist or materials scientist hypothesized a composition based on known crystal chemistry, synthesized a small sample, characterized its structure and relevant properties through diffraction, spectroscopy, and electrical measurement, and evaluated whether the result warranted further investigation. Each iteration consumed days to weeks of laboratory time. The cost of identifying a single promising new compound through this process, including failed attempts, ran to tens or hundreds of thousands of dollars in personnel and equipment.

High-throughput computational methods beginning in the 2000s partially automated the screening stage. Density functional theory (DFT) calculations can predict the thermodynamic stability and some electronic properties of a hypothetical compound without synthesizing it, at a computational cost that fell dramatically as processor performance improved. The Materials Project, a federally funded database launched at Lawrence Berkeley National Laboratory in 2011, grew to contain computational predictions for more than 150,000 inorganic compounds by the early 2020s, each characterized for stability and

key electronic properties.¹ This database became the training substrate for the machine learning approaches that followed.

What Graph Networks for Materials Exploration Demonstrated

In November 2023, researchers at Google DeepMind published results from GNoME – Graph Networks for Materials Exploration – a deep learning system trained on Materials Project data and extended through active learning.² The system predicted 2.2 million new inorganic crystal structures, of which 381,000 were assessed as thermodynamically stable. This expanded the universe of computationally predicted stable materials by roughly a factor of ten. DeepMind made the 381,000 stable structures publicly available, and in concurrent work, researchers at Lawrence Berkeley National Laboratory demonstrated that a robotic synthesis facility called A-Lab could use the predictions to guide autonomous experimental synthesis efforts, producing 41 of 58 targeted novel compounds within a 17-day campaign.³

The results are genuine advances. The scale of screening that GNoME accomplished would have required centuries of DFT computation at pre-machine-learning rates. The A-Lab demonstration showed that the loop between computational prediction and experimental realization can be at least partially closed through automation. For specific material families – particularly layered compounds with two-dimensional character, of which GNoME identified roughly 52,000 candidate structures – the predictions point toward a large unexplored space that could yield new materials for semiconductor applications, superconductivity, and battery electrolytes.

The GNoME results also require careful reading. The more fundamental limitation is economic rather than methodological: thermodynamic stability, which GNoME predicts, is a necessary but not sufficient condition for a material to be useful in manufacturing.

The Synthesizability Problem

A crystal structure that a deep learning model predicts to be thermodynamically stable at zero Kelvin and ambient pressure may be impossible to synthesize under practical laboratory or industrial conditions for several reasons.

The synthesis pathway matters. Thermodynamics determines whether a phase is stable at equilibrium; kinetics determines whether a synthesis process can reach that phase within a practical timeframe. Many thermodynamically stable compounds are kinetically inaccessible – the activation barriers to forming the desired phase are too high under conditions that do not also destroy the product. Machine learning models trained on

stability data do not inherently encode kinetic accessibility.

Precursor availability and process compatibility matter. Industrial materials processes are not arbitrary; they use specific reagents, temperatures, and atmospheres that are compatible with the surrounding process environment. A semiconductor dielectric material, for example, must not only have suitable electrical properties but must also be depositable using chemical vapor deposition or atomic layer deposition chemistry that integrates with existing tool sets and does not contaminate adjacent materials. Predicting that a material has the right bandgap does not predict whether it can be deposited conformally at 300 degrees Celsius from available precursor gases.

Phase stability under processing conditions matters. Many semiconductor processes subject materials to high temperatures, reactive plasmas, and aggressive wet chemistries. A material stable at standard conditions may decompose, crystallize, or interdiffuse with neighboring layers under process conditions. Predicting ambient stability does not predict process stability.

The A-Lab synthesis success rate – 41 of 58 targeted compounds, roughly 70 percent – illustrates the gap even in a carefully controlled academic setting, where targets were presumably selected partly for synthesis accessibility. In industrial semiconductor development, the synthesis success rate for computationally predicted candidate process materials is considerably lower, and the cost of a failed candidate that passes through months of process integration before failing is much higher.

Where Machine Learning Generates Durable Economic Value

The limitations above do not undermine the value of machine learning in materials discovery; they define where that value is most concentrated.

Property prediction on known synthesis routes is the highest-value near-term application. If a researcher knows how to make a class of materials – for example, transition metal dichalcogenides deposited by atomic layer deposition – machine learning models can screen thousands of compositional variants to identify which are most likely to exhibit a target property, dramatically reducing the experimental effort required. This use case requires only that the model generalize within a known synthesis space, not that it predict synthesizability from scratch.

Process parameter optimization is a second high-value application. Thin-film deposition processes such as ALD involve tens of adjustable parameters: precursor dose times, purge durations, temperature, pressure, and plasma conditions. Finding the combination of

parameters that achieves target film properties – thickness uniformity, impurity content, interfacial chemistry – through experiment alone requires large numbers of depositions. Machine learning models trained on process data can identify promising parameter combinations with fewer experiments, shortening the development timeline for new processes. This application does not require predicting new materials; it requires optimizing within a known process space.

Defect characterization and root cause analysis in manufacturing is a third application, now in active industrial deployment. Wafer defect patterns encode information about which process step and which equipment generated the defect. Machine learning classifiers trained on historical defect images and process logs can identify defect signatures and trace them to specific upstream causes faster than manual analysis, reducing the time between defect detection and process correction.

Economic Implications

The economic significance of machine learning for materials discovery is therefore not evenly distributed across the discovery pipeline. At the hypothesis and screening stage, AI tools create large value by expanding the space of candidates that can be evaluated computationally. At the synthesis and process integration stage, the value is real but more bounded, because the synthesis bottleneck that AI has not yet solved limits how quickly computational discoveries can translate into manufacturable materials.

For semiconductor process development specifically, where new materials must integrate with existing tool sets and process flows that have been refined over decades, the path from a GNoME-predicted crystal structure to a deployed process layer is measured in years of engineering work even when the computational prediction is correct. The economic implication is that machine learning accelerates the early discovery pipeline substantially while leaving the experimental validation and process integration pipeline largely unchanged. Firms that understand this distinction will calibrate their investment in computational materials discovery accordingly, concentrating resources where the screening bottleneck is genuinely rate-limiting and maintaining the experimental infrastructure needed to evaluate candidates the computational pipeline generates.

References

¹ Anubhav Jain et al., “Commentary: The Materials Project: A Materials Genome Approach to Accelerating Materials Innovation,” *APL Materials* 1, no. 1 (2013): 011002.

² Amil Merchant et al., “Scaling Deep Learning for Materials Discovery,” *Nature* 624 (November 2023): 80-85.

³ Bernardus Reniers et al., “Autonomous Chemical Research with Large Language Models,” *Nature* 624 (November 2023): 570-578 (the A-Lab work, published concurrently with GNoME).

⁴ This point is acknowledged in the original GNoME paper itself. Merchant et al. note that their stability predictions rely on DFT calculations at standard conditions and that kinetic accessibility, synthesizability under practical conditions, and compatibility with existing process environments are not captured by the thermodynamic stability metric. The A-Lab synthesis campaign’s 71 percent success rate on a curated target list (41 of 58 compounds) is consistent with this limitation even under favorable conditions.